

Class 8 Breast Cancer Mini Project

Patrick Nguyen (ID:A17680785)

Table of contents

Background	1
Data Import	1
Principal Component Analysis (PCA)	3
Scree plot	7
4. Hierarchical clustering	9
Combining methods	10
Prediction	13

Background

In today's class we will be employing all the R techniques for data analysis that we have learned thus far - including the machine learning methods of clustering and PCA - to analyze real breast cancer biopsy data.

Data Import

The data is in CSV format:

```
#read.csv("WisconsinCancer.csv")  
  
wisc.df <- read.csv("WisconsinCancer.csv", row.names=1)  
#wisc.df
```

wee peak at the data

```
head(wisc.df, 3)
```

	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean
842302	M	17.99	10.38	122.8	1001
842517	M	20.57	17.77	132.9	1326
84300903	M	19.69	21.25	130.0	1203
	smoothness_mean	compactness_mean	concavity_mean	concave.points_mean	
842302	0.11840	0.27760	0.3001	0.14710	
842517	0.08474	0.07864	0.0869	0.07017	
84300903	0.10960	0.15990	0.1974	0.12790	
	symmetry_mean	fractal_dimension_mean	radius_se	texture_se	perimeter_se
842302	0.2419	0.07871	1.0950	0.9053	8.589
842517	0.1812	0.05667	0.5435	0.7339	3.398
84300903	0.2069	0.05999	0.7456	0.7869	4.585
	area_se	smoothness_se	compactness_se	concavity_se	concave.points_se
842302	153.40	0.006399	0.04904	0.05373	0.01587
842517	74.08	0.005225	0.01308	0.01860	0.01340
84300903	94.03	0.006150	0.04006	0.03832	0.02058
	symmetry_se	fractal_dimension_se	radius_worst	texture_worst	
842302	0.03003	0.006193	25.38	17.33	
842517	0.01389	0.003532	24.99	23.41	
84300903	0.02250	0.004571	23.57	25.53	
	perimeter_worst	area_worst	smoothness_worst	compactness_worst	
842302	184.6	2019	0.1622	0.6656	
842517	158.8	1956	0.1238	0.1866	
84300903	152.5	1709	0.1444	0.4245	
	concavity_worst	concave.points_worst	symmetry_worst		
842302	0.7119	0.2654	0.4601		
842517	0.2416	0.1860	0.2750		
84300903	0.4504	0.2430	0.3613		
	fractal_dimension_worst				
842302	0.11890				
842517	0.08902				
84300903	0.08758				

Q1. How many observations are in this dataset?

```
nrow(wisc.df)
```

[1] 569

Q2. How many of the observations have a malignant diagnosis?

```
sum(wisc.df$diagnosis == "M")
```

```
[1] 212
```

```
table(wisc.df$diagnosis)
```

```
B      M  
357 212
```

Q3. How many variables/features in the data are suffixed with _mean?

```
length(grep("_mean", colnames(wisc.df)))
```

```
[1] 10
```

We need to remove the diagnosis column before we do any further analysis of this dataset - we don't want to pass this to PCA etc. We will save it as a separate vector that we can use later to compare our findings to those of experts.

```
wisc.data <- wisc.df[ , -1]  
diagnosis <- wisc.df$diagnosis
```

Principal Component Analysis (PCA)

The main function in base R is called `prcomp()` we will use the optional argument `scale=TRUE` here as the data columns/features.dimensions are on very different scales in the original data set.


```
wisc.pr <- prcomp(wisc.data, scale=T)
```

```
attributes(wisc.pr)
```

```
$names  
[1] "sdev"      "rotation"   "center"    "scale"     "x"  
  
$class  
[1] "prcomp"
```

```
library(ggplot2)

ggplot(wisc.pr$x) +
  aes(PC1, PC2, col=diagnosis) +
  geom_point()
```



Q4. From your results, what proportion of the original variance is captured by the first principal component (PC1)?

```
summary(wisc.pr)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	3.6444	2.3857	1.67867	1.40735	1.28403	1.09880	0.82172
Proportion of Variance	0.4427	0.1897	0.09393	0.06602	0.05496	0.04025	0.02251
Cumulative Proportion	0.4427	0.6324	0.72636	0.79239	0.84734	0.88759	0.91010
	PC8	PC9	PC10	PC11	PC12	PC13	PC14
Standard deviation	0.69037	0.6457	0.59219	0.5421	0.51104	0.49128	0.39624
Proportion of Variance	0.01589	0.0139	0.01169	0.0098	0.00871	0.00805	0.00523
Cumulative Proportion	0.92598	0.9399	0.95157	0.9614	0.97007	0.97812	0.98335
	PC15	PC16	PC17	PC18	PC19	PC20	PC21
Standard deviation	0.30681	0.28260	0.24372	0.22939	0.22244	0.17652	0.1731

```

Proportion of Variance 0.00314 0.00266 0.00198 0.00175 0.00165 0.00104 0.0010
Cumulative Proportion 0.98649 0.98915 0.99113 0.99288 0.99453 0.99557 0.9966
PC22 PC23 PC24 PC25 PC26 PC27 PC28
Standard deviation 0.16565 0.15602 0.1344 0.12442 0.09043 0.08307 0.03987
Proportion of Variance 0.00091 0.00081 0.0006 0.00052 0.00027 0.00023 0.00005
Cumulative Proportion 0.99749 0.99830 0.9989 0.99942 0.99969 0.99992 0.99997
PC29 PC30
Standard deviation 0.02736 0.01153
Proportion of Variance 0.00002 0.00000
Cumulative Proportion 1.00000 1.00000

```

44.27% of variance was captured

Q5. How many principal components (PCs) are required to describe at least 70% of the original variance in the data?

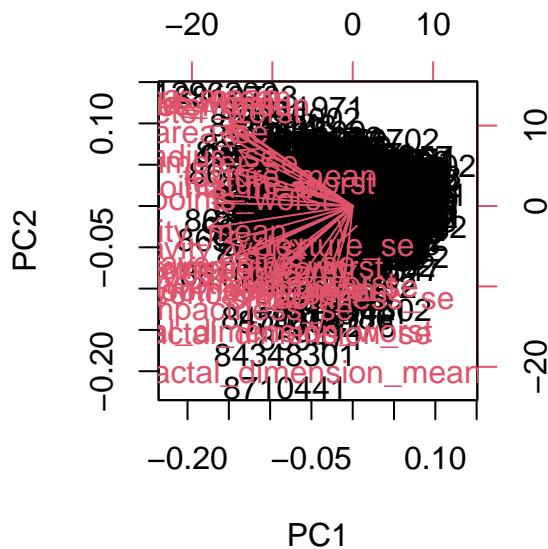
3 principal components (PCs 1 to 3) are needed to describe 70% of the original variance

Q6. How many principal components (PCs) are required to describe at least 90% of the original variance in the data?

7 principal components (PCs 1 to 7) are needed to describe 90% of the original variance

Q7. What stands out to you about this plot? Is it easy or difficult to understand?
Why?

```
biplot(wisc.pr)
```

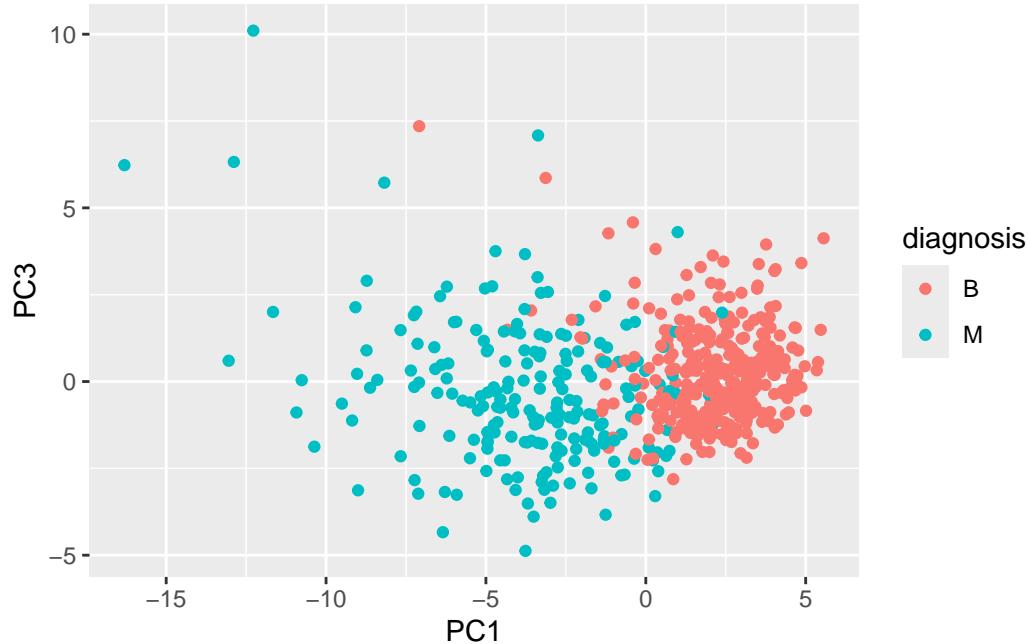


This plot is hard to read due to the large amount of labels on the graph making it difficult to discern what each label means as well as seeing the points on the graph

Q8. Generate a similar plot for principal components 1 and 3. What do you notice about these plots?

```
library(ggplot2)

ggplot(wisc.pr$x) +
  aes(PC1, PC3, col=diagnosis) +
  geom_point()
```



this scatter plot has a similar spread to the PC1 vs PC2 plot but this plot has the points located generally at a lower y-axis area than the PC1 vs PC2 plot.


Scree plot

```
library(factoextra)
```

Welcome! Want to learn more? See two factoextra-related books at <https://goo.gl/ve3WBa>

```
fviz_eig(wisc.pr, addlabels = TRUE)
```

Warning in geom_bar(stat = "identity", fill = barfill, color = barcolor, :
Ignoring empty aesthetic: `width`.



Q9. For the first principal component, what is the component of the loading vector (i.e. `wisc.pr$rotation[,1]`) for the feature `concave.points_mean`? This tells us how much this original feature contributes to the first PC. Are there any features with larger contributions than this one?

```
pc1_loading_concave <- wisc.pr$rotation["concave.points_mean", "PC1"]

pc1_loading_concave
```

[1] -0.2608538

```
pc1_loadings <- wisc.pr$rotation[, "PC1"]

ordered_loadings <- sort(abs(pc1_loadings), decreasing = TRUE)

head(ordered_loadings, 10)
```

concave.points_mean	0.2608538	concavity_mean	concave.points_worst
compactness_mean		0.2584005	0.2508860
		perimeter_worst	concavity_worst

```
0.2392854          0.2366397          0.2287675  
radius_worst        perimeter_mean      area_worst  
0.2279966          0.2275373          0.2248705  
area_mean             
0.2209950
```

```
abs(pc1_loading_concave)
```

```
[1] 0.2608538
```

```
ordered_loadings["concave.points_mean"]
```

```
concave.points_mean  
0.2608538
```

concave.points_mean is the largest contributor, ever other contributor is smaller than concave.points_mean

4. Hierarchical clustering

The goal of this section is to do hierarchical clustering of the original data to see if there is any obvious grouping into malignant and benign clusters

In short, these results are not good!

First we will scale our `wisc.data` then calculate a distance matrix

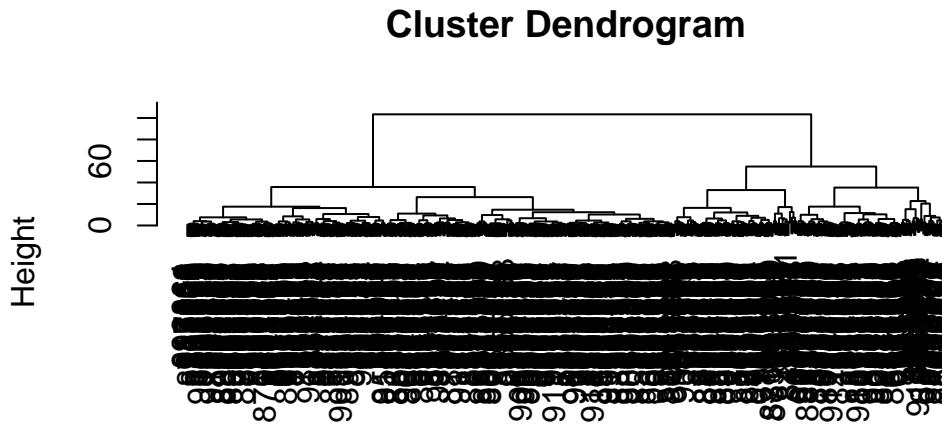
```
wisc.hclust <- hclust( dist( scale(wisc.data)))  
wisc.hclust.clusters <- cutree(wisc.hclust, k=2)  
table(wisc.hclust.clusters)
```

```
wisc.hclust.clusters  
1    2  
567  2
```

Combining methods

The idea here is that I can take my new variables (i.e. the scores on the PCs `wisc.pr$x`) that are better descriptors of the data-set than the original features (i.e. the 30 columns in `wisc.data`) and use these as a basis for clustering.

```
pc.dist <- dist(wisc.pr$x[, 1:3])
wisc.pr.hclust <- hclust(pc.dist, method = "ward.D2")
plot(wisc.pr.hclust)
```



```
pc.dist
hclust (*, "ward.D2")
```

```
grps <- cutree(wisc.pr.hclust, k=2)
table(grps)
```

```
grps
 1   2
203 366
```

I can now run `table()` with both my clustering `grps` and the expert diagnosis

```
table(grps, diagnosis)
```

```
diagnosis
grps   B    M
1    24  179
2   333   33
```

Our cluster “1” has 179 “M” diagnosis or cluster “2” has 333 “B” diagnosis

179 TP 24 FP 333 TN 33 FN

Sensitivity: TP/(TP+FN)

```
179/(179+33)
```

```
[1] 0.8443396
```

Specificity: TN/(TN+FP)


```
333/(333+24)
```

```
[1] 0.9327731
```

Q10. Using the plot() and abline() functions, what is the height at which the clustering model has 4 clusters?

```
wisc.hclust <- hclust(dist(scale(wisc.data)))
plot(wisc.hclust)
abline(h = 4, col="red", lty=2)
```

Cluster Dendrogram



```
dist(scale(wisc.data))
hclust (*, "complete")
```

at around a height of 15 is when the cluster model has 4 clusters.

Q12. Which method gives your favorite results for the same data.dist dataset?
Explain your reasoning.

I prefer ward.d2 because in my opinion it organizes the clusters the best out of all the methods which makes interpreting the cluster model not too difficult compared to other methods as the other methods make the clusters look too cluttered.

Q14. How well do the hierarchical clustering models you created in the previous sections (i.e. without first doing PCA) do in terms of separating the diagnoses? Again, use the table() function to compare the output of each model (wisc.hclust.clusters and wisc.pr.hclust.clusters) with the vector containing the actual diagnoses.

```
table(wisc.hclust.clusters, diagnosis)
```

wisc.hclust.clusters	B	M
1	357	210
2	0	2

```
wisc.pr.hclust.clusters <- cutree(wisc.pr.hclust, k = 2)
table(wisc.pr.hclust.clusters, diagnosis)
```

```

diagnosis
wisc.pr.hclust.clusters   B    M
      1  24 179
      2 333 33

```

The “PR” cluster separates better than the other cluster model as the ratios of Benign and Malignant for the “PR” cluster model are larger than the other cluster model.

Prediction

We can use our PCA model for prediction of new un-seen cases.

```

#url <- "new_samples.csv"
url <- "https://tinyurl.com/new-samples-CSV"
new <- read.csv(url)
npc <- predict(wisc.pr, newdata=new)
npc


```

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
[1,]	2.576616	-3.135913	1.3990492	-0.7631950	2.781648	-0.8150185	-0.3959098
[2,]	-4.754928	-3.009033	-0.1660946	-0.6052952	-1.140698	-1.2189945	0.8193031
	PC8	PC9	PC10	PC11	PC12	PC13	PC14
[1,]	-0.2307350	0.1029569	-0.9272861	0.3411457	0.375921	0.1610764	1.187882
[2,]	-0.3307423	0.5281896	-0.4855301	0.7173233	-1.185917	0.5893856	0.303029
	PC15	PC16	PC17	PC18	PC19	PC20	
[1,]	0.3216974	-0.1743616	-0.07875393	-0.11207028	-0.08802955	-0.2495216	
[2,]	0.1299153	0.1448061	-0.40509706	0.06565549	0.25591230	-0.4289500	
	PC21	PC22	PC23	PC24	PC25	PC26	
[1,]	0.1228233	0.09358453	0.08347651	0.1223396	0.02124121	0.078884581	
[2,]	-0.1224776	0.01732146	0.06316631	-0.2338618	-0.20755948	-0.009833238	
	PC27	PC28	PC29	PC30			
[1,]	0.220199544	-0.02946023	-0.015620933	0.005269029			
[2,]	-0.001134152	0.09638361	0.002795349	-0.019015820			

```

plot(wisc.pr$x[,1:2], col=grps)
points(npc[,1], npc[,2], col="blue", pch=16, cex=3)
text(npc[,1], npc[,2], c(1,2), col="white")

```



Q16. Which of these new patients should we prioritize for follow up based on your results?

Based on the graph the 2nd group of patients (2 on graph) should be prioritized due to the high amount of patients that have malignant cancer.