

class 9 candy

Patrick Nguyen (ID: A17680785)

Table of contents

Background	1
Data Import	1
Exploratory Analysis	4
Overall Candy Rankings	7
Taking a look at pricepercent	16
Exploring the correlation structure	17
Principal Component Analysis	18

Background

In today's mini-project we will analyze candy data with exploratory graphics , basic statistics , correlation analysis and principal component analysis methods we have been learning thus far.

Data Import

The data comes as a CSV file from 538.

```
candy_file <- ("candy-data.csv")
candy = read.csv(candy_file, row.names=1)
head(candy)
```

	chocolate	fruity	caramel	peanut	almond	nougat	crisped	rice	wafer
100 Grand	1	0	1			0	0		1
3 Musketeers	1	0	0			0	1		0
One dime	0	0	0			0	0		0
One quarter	0	0	0			0	0		0

Air Heads	0	1	0	0	0	0
Almond Joy	1	0	0	1	0	0
hard bar pluribus sugarpercent pricepercent winpercent						
100 Grand	0	1	0	0.732	0.860	66.97173
3 Musketeers	0	1	0	0.604	0.511	67.60294
One dime	0	0	0	0.011	0.116	32.26109
One quarter	0	0	0	0.011	0.511	46.11650
Air Heads	0	0	0	0.906	0.511	52.34146
Almond Joy	0	1	0	0.465	0.767	50.34755

Q1. How many different candy types are in this dataset?

```
nrow(candy)
```

[1] 85

Q2. How many fruity candy types are in the dataset?

```
sum(candy$fruity)
```

[1] 38

Q3. What is your favorite candy (other than Twix) in the dataset and what is its winpercent value?

```
candy["Nestle Crunch", ]$winpercent
```

[1] 66.47068

Q4. What is the winpercent value for “Kit Kat”?

```
candy["Kit Kat", ]$winpercent
```

[1] 76.7686

Q5. What is the winpercent value for “Tootsie Roll Snack Bars”?

```
candy["Tootsie Roll Snack Bars", ]$winpercent
```

[1] 49.6535

Q6. Is there any variable/column that looks to be on a different scale to the majority of the other columns in the dataset?

```
library("skimr")
skim(candy)
```

Table 1: Data summary

Name	candy
Number of rows	85
Number of columns	12
Column type frequency:	
numeric	12
Group variables	None

Variable type: numeric

skim_variable	n_missing	com-	mean	sd	p0	p25	p50	p75	p100	hist
		plete_rate								
chocolate	0	1	0.44	0.50	0.00	0.00	0.00	1.00	1.00	
fruity	0	1	0.45	0.50	0.00	0.00	0.00	1.00	1.00	
caramel	0	1	0.16	0.37	0.00	0.00	0.00	0.00	1.00	
peanutyaly-	0	1	0.16	0.37	0.00	0.00	0.00	0.00	1.00	
mondy										
nougat	0	1	0.08	0.28	0.00	0.00	0.00	0.00	1.00	
crispedrice-	0	1	0.08	0.28	0.00	0.00	0.00	0.00	1.00	
wafer										
hard	0	1	0.18	0.38	0.00	0.00	0.00	0.00	1.00	
bar	0	1	0.25	0.43	0.00	0.00	0.00	0.00	1.00	
pluribus	0	1	0.52	0.50	0.00	0.00	1.00	1.00	1.00	
sugarpercent	0	1	0.48	0.28	0.01	0.22	0.47	0.73	0.99	
pricepercent	0	1	0.47	0.29	0.01	0.26	0.47	0.65	0.98	
winpercent	0	1	50.32	14.71	22.45	39.14	47.83	59.86	84.18	

The “winpercent” variable looks like it is on a different scale because the numbers it has for each column are a lot larger than the other variables.

Q7. What do you think a zero and one represent for the candy\$chocolate column?

```
candy$chocolate
```

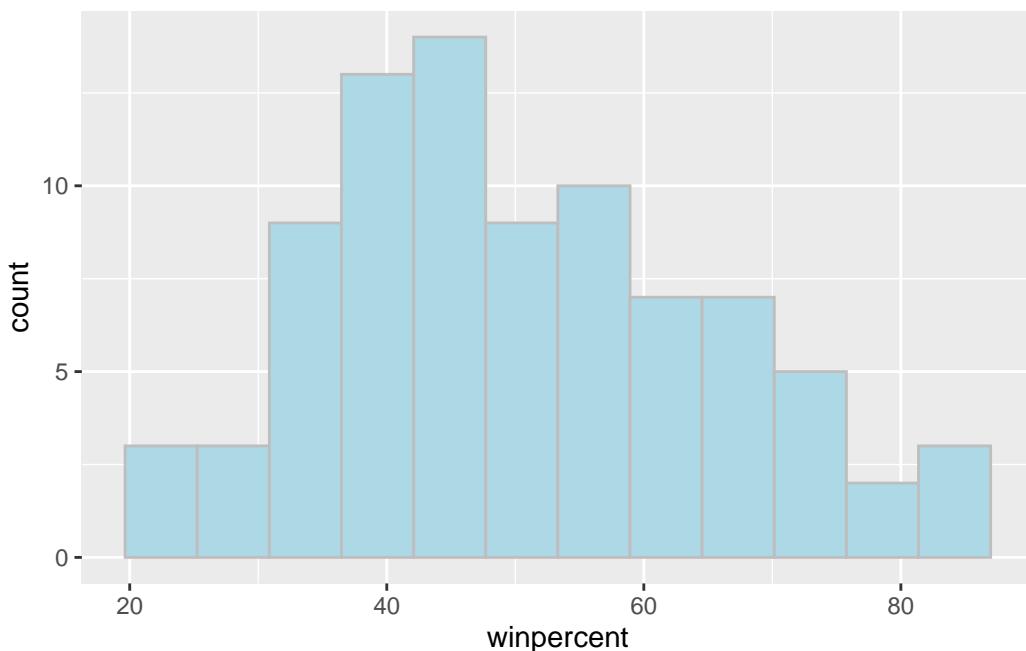
```
[1] 1 1 0 0 0 1 1 0 0 0 1 0 0 0 0 0 0 0 0 0 0 1 1 1 0 1 1 1 0 0 0 1 1 0 1 1 1  
[39] 1 1 1 0 1 1 0 0 0 1 0 0 0 1 1 1 1 0 1 0 0 1 0 0 1 0 1 1 0 0 0 0 0 0 0 1 1  
[77] 1 1 0 1 0 0 0 0 1
```

The “0” likely represents a candy that doesn’t have any chocolate while the “1” represents candy that does have chocolate in it.

Exploratory Analysis

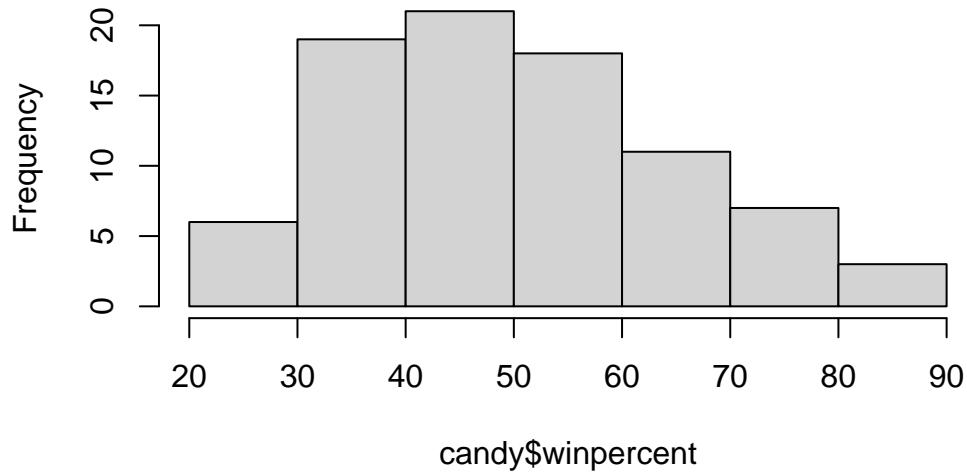
Q8. Plot a histogram of winpercent values using both base R an ggplot2.

```
library(ggplot2)
ggplot(candy) +
  aes(winpercent) +
  geom_histogram(bins=12, fill="lightblue", col="gray")
```



```
hist(candy$winpercent, breaks=8)
```

Histogram of candy\$winpercent



Q9. Is the distribution of winpercent values symmetrical?

The distribution is not symmetrical

Q10. Is the center of the distribution above or below 50%?

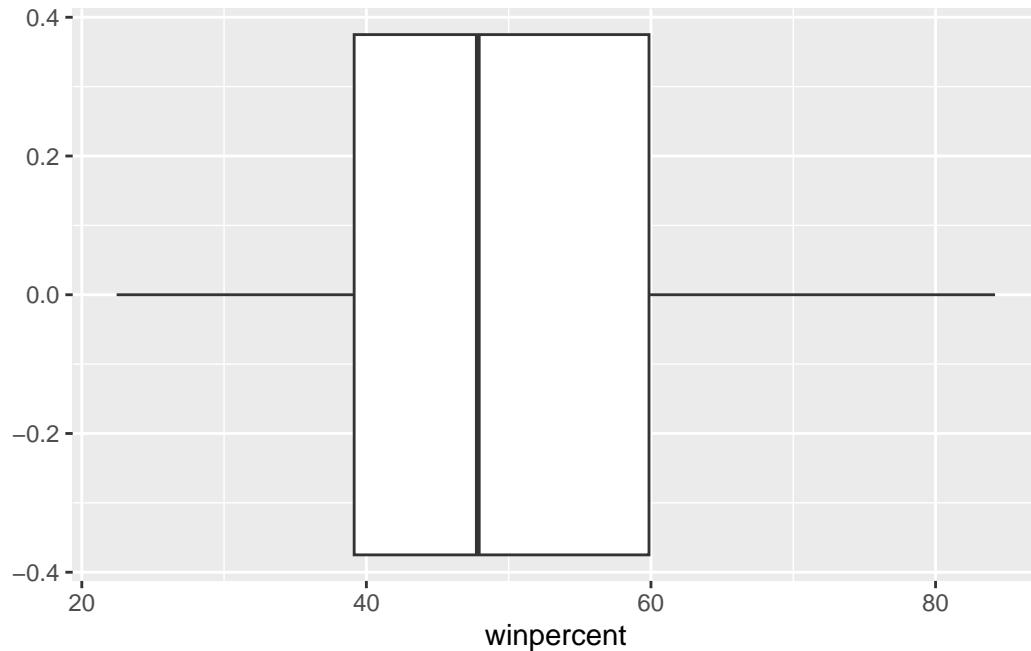
```
mean(candy$winpercent)
```

[1] 50.31676

```
summary(candy$winpercent)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
22.45	39.14	47.83	50.32	59.86	84.18

```
ggplot(candy) +  
  aes(winpercent) +  
  geom_boxplot()
```



The center of the distribution is below 50%

Q11. On average is chocolate candy higher or lower ranked than fruit candy?

Steps to solve this: 1.Find all chocolate candy in the dataset 2.Extract or find their winpercent values 3.Calculate the mean of these values

4.Find all fruit candy in the data set 5.Find their winpercent 6.Calculate their mean value

```
chocolate_can <- candy[ candy$chocolate==1, ]
choc.win <- chocolate_can$winpercent
mean(choc.win)
```

[1] 60.92153

```
fruit_can <- candy[ candy$fruity==1, ]
fruit.win <- fruit_can$winpercent
mean(fruit.win)
```

[1] 44.11974

chocolate candy on average is higher ranked than fruit candy.

Q12. Is this difference statistically significant?

```
t.test(choc.win, fruit.win)
```

Welch Two Sample t-test

```
data: choc.win and fruit.win
t = 6.2582, df = 68.882, p-value = 2.871e-08
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 11.44563 22.15795
sample estimates:
mean of x mean of y
 60.92153 44.11974
```

The difference between these two candies is statistically significant

Overall Candy Rankings

Q13. What are the five least liked candy types in this set?

```
head(candy[order(candy$winpercent),], n=5)
```

	chocolate	fruity	caramel	peanut	almond	nougat	
Nik L Nip	0	1	0	0	0	0	
Boston Baked Beans	0	0	0	1	0	0	
Chiclets	0	1	0	0	0	0	
Super Bubble	0	1	0	0	0	0	
Jawbusters	0	1	0	0	0	0	
	crispedrice	wafers	hard	bar	pluribus	sugarpercent	pricepercent
Nik L Nip	0	0	0	1	0.197	0.976	
Boston Baked Beans	0	0	0	1	0.313	0.511	
Chiclets	0	0	0	1	0.046	0.325	
Super Bubble	0	0	0	0	0.162	0.116	
Jawbusters	0	1	0	1	0.093	0.511	
	winpercent						
Nik L Nip	22.44534						
Boston Baked Beans	23.41782						
Chiclets	24.52499						
Super Bubble	27.30386						
Jawbusters	28.12744						

Q14. What are the top 5 all time favorite candy types out of this set?

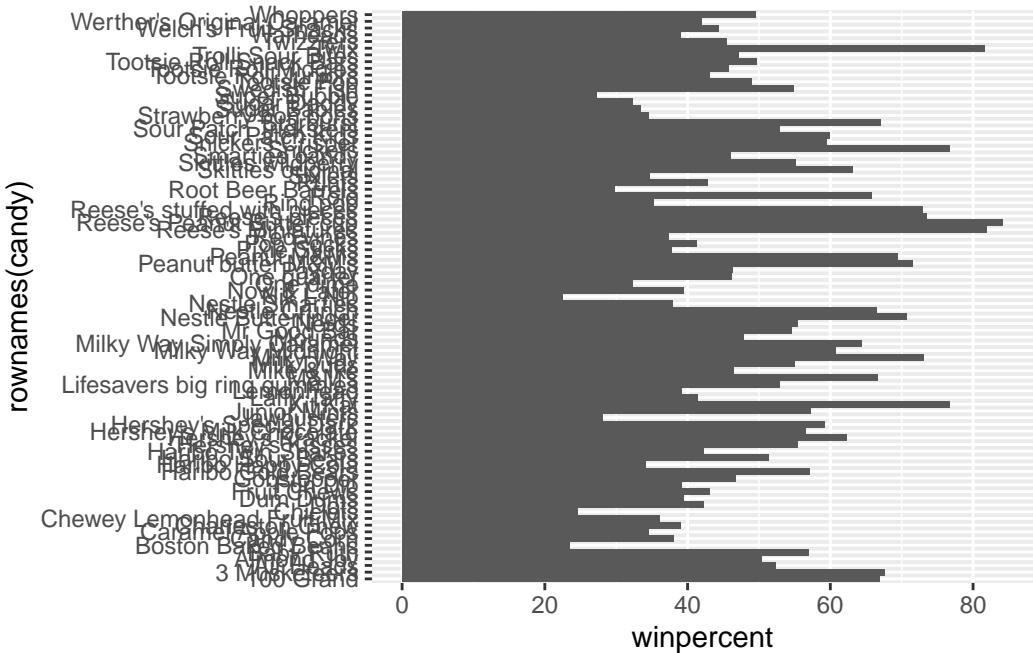
```
tail(candy[order(candy$winpercent),], n=5)
```

	chocolate	fruity	caramel	peanut	yalmond	nougat
Snickers	1	0	1		1	1
Kit Kat	1	0	0		0	0
Twix	1	0	1		0	0
Reese's Miniatures	1	0	0		1	0
Reese's Peanut Butter cup	1	0	0		1	0
	crispedrice	wafers	hard bar	pluribus	sugar	percent
Snickers	0	0	1	0	0.546	
Kit Kat	1	0	1	0	0.313	
Twix	1	0	1	0	0.546	
Reese's Miniatures	0	0	0	0	0.034	
Reese's Peanut Butter cup	0	0	0	0	0.720	
	price	percent	winpercent			
Snickers	0.651	76.67378				
Kit Kat	0.511	76.76860				
Twix	0.906	81.64291				
Reese's Miniatures	0.279	81.86626				
Reese's Peanut Butter cup	0.651	84.18029				

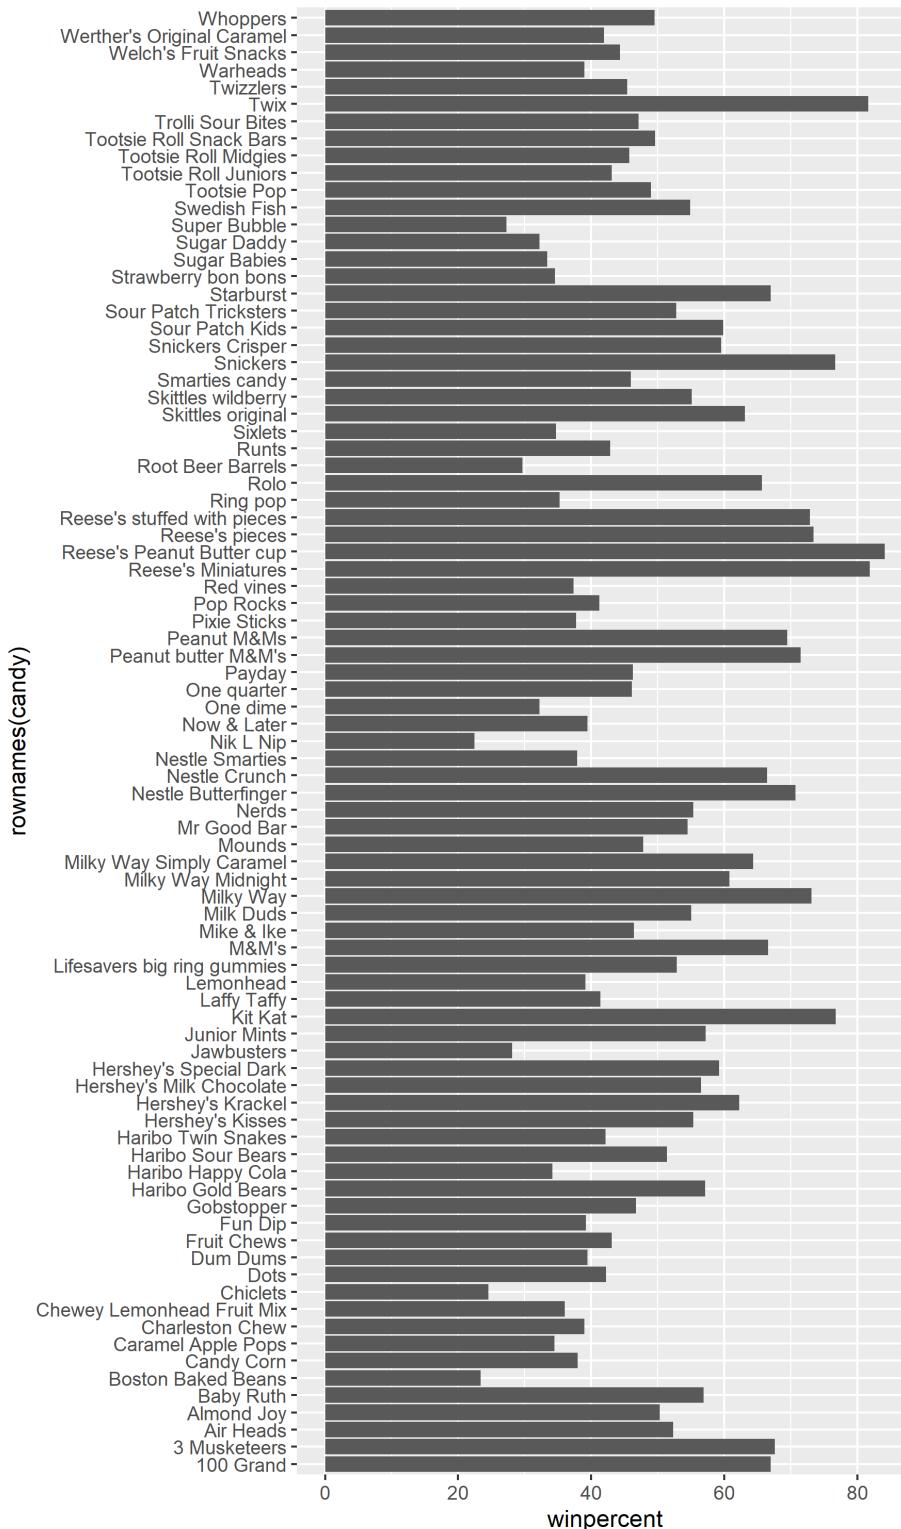
Q15. Make a first barplot of candy ranking based on winpercent values.

```
library(ggplot2)

ggplot(candy) +
  aes(winpercent, rownames(candy)) +
  geom_col()
```

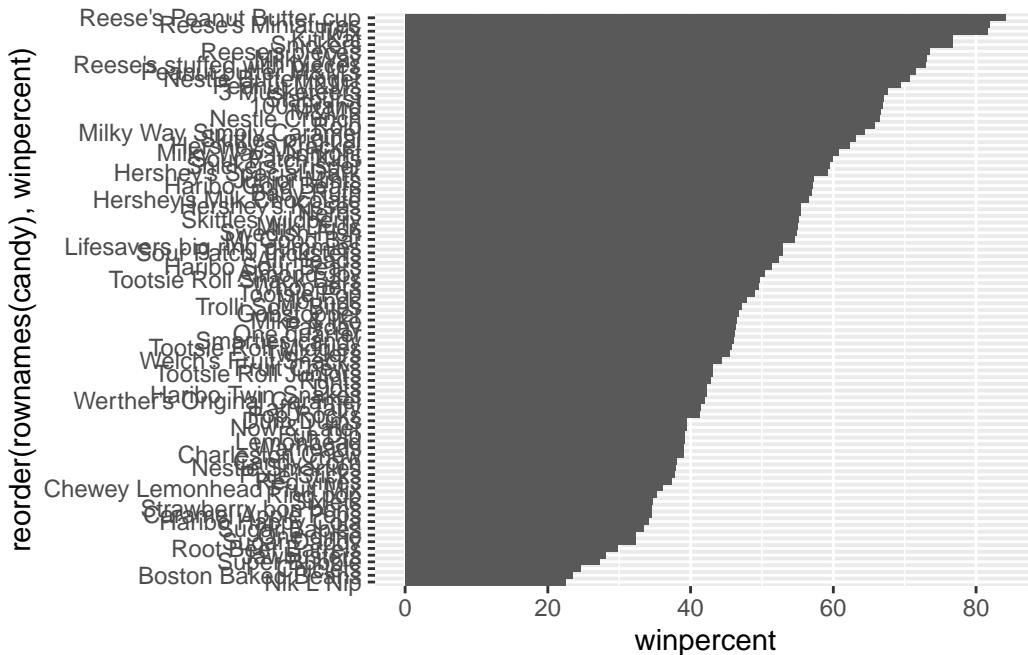


```
ggsave("barplot1.png", height=10, width=6)
```

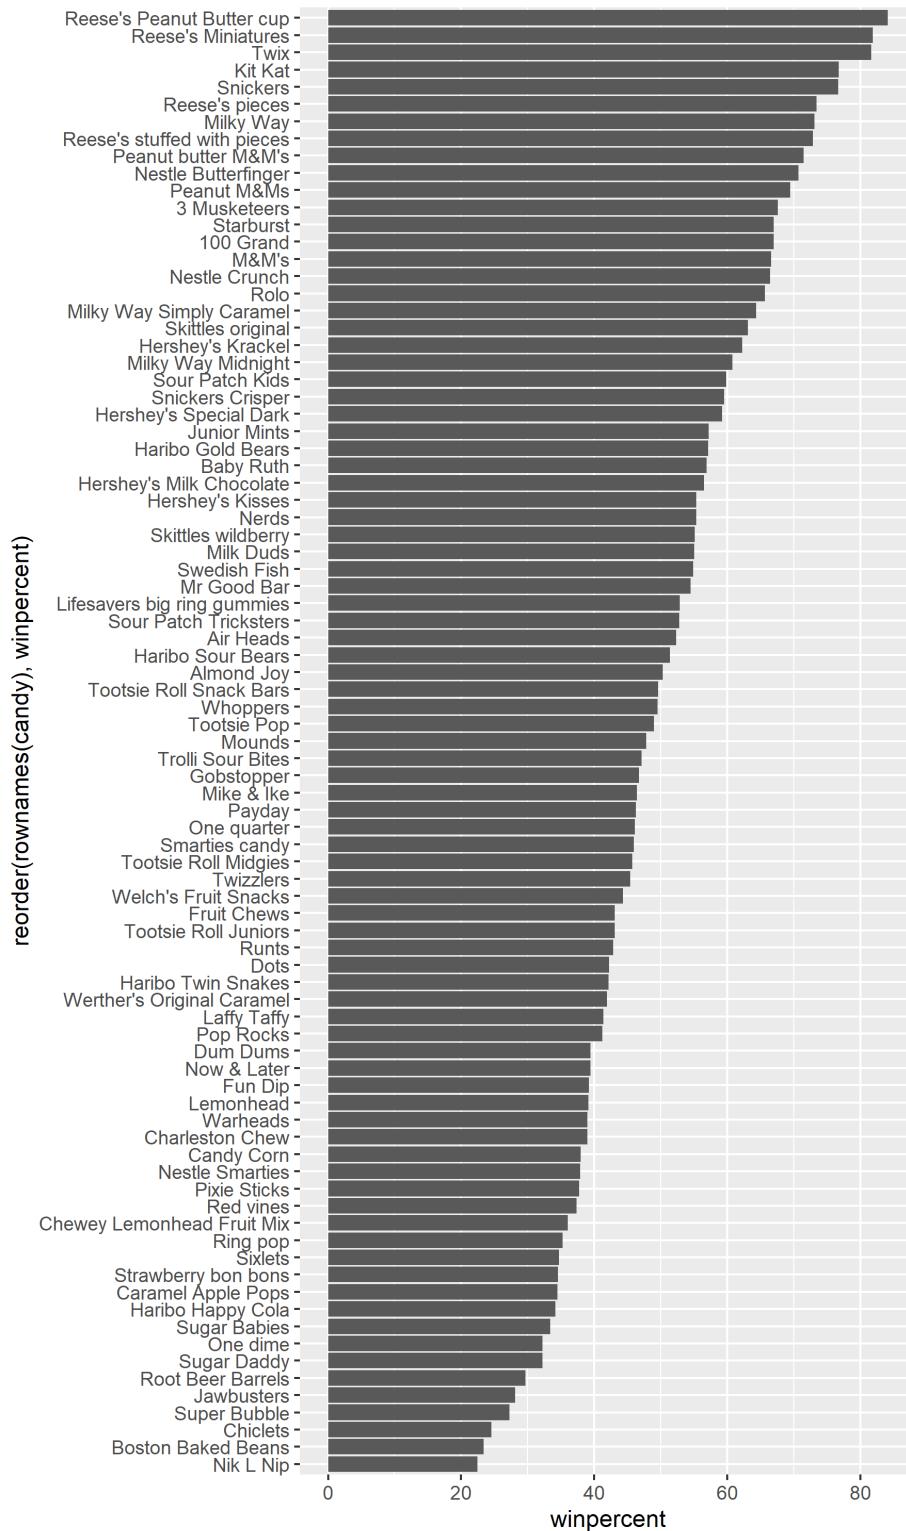


Q16. This is quite ugly, use the reorder() function to get the bars sorted by winpercent?

```
ggplot(candy) +  
  aes(winpercent, reorder(rownames(candy),winpercent)) +  
  geom_col()
```



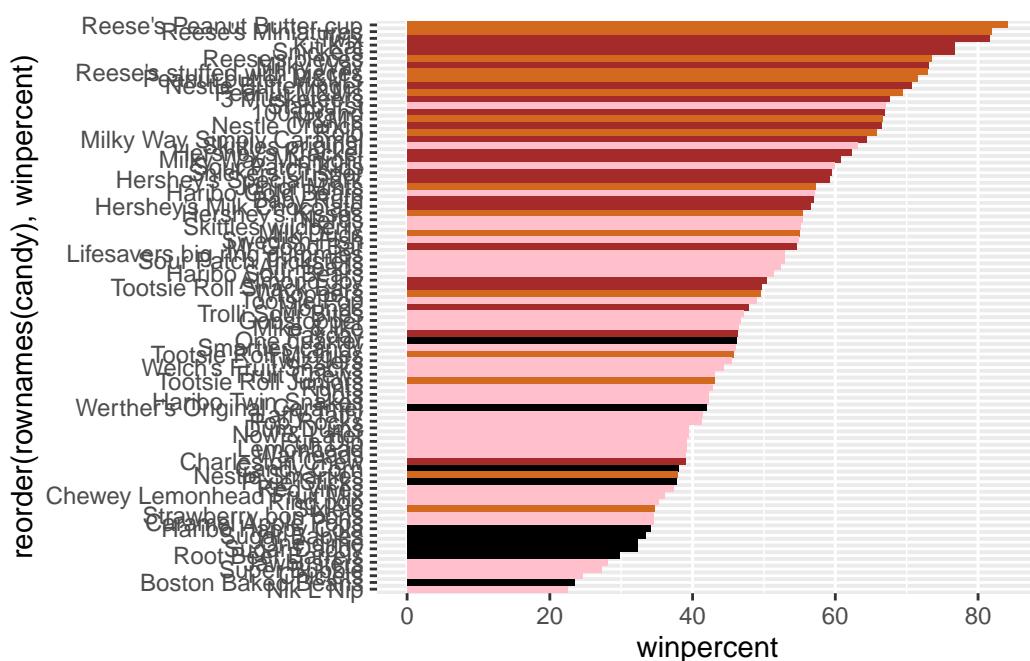
```
ggsave("barplot2.png", height=10, width=6)
```



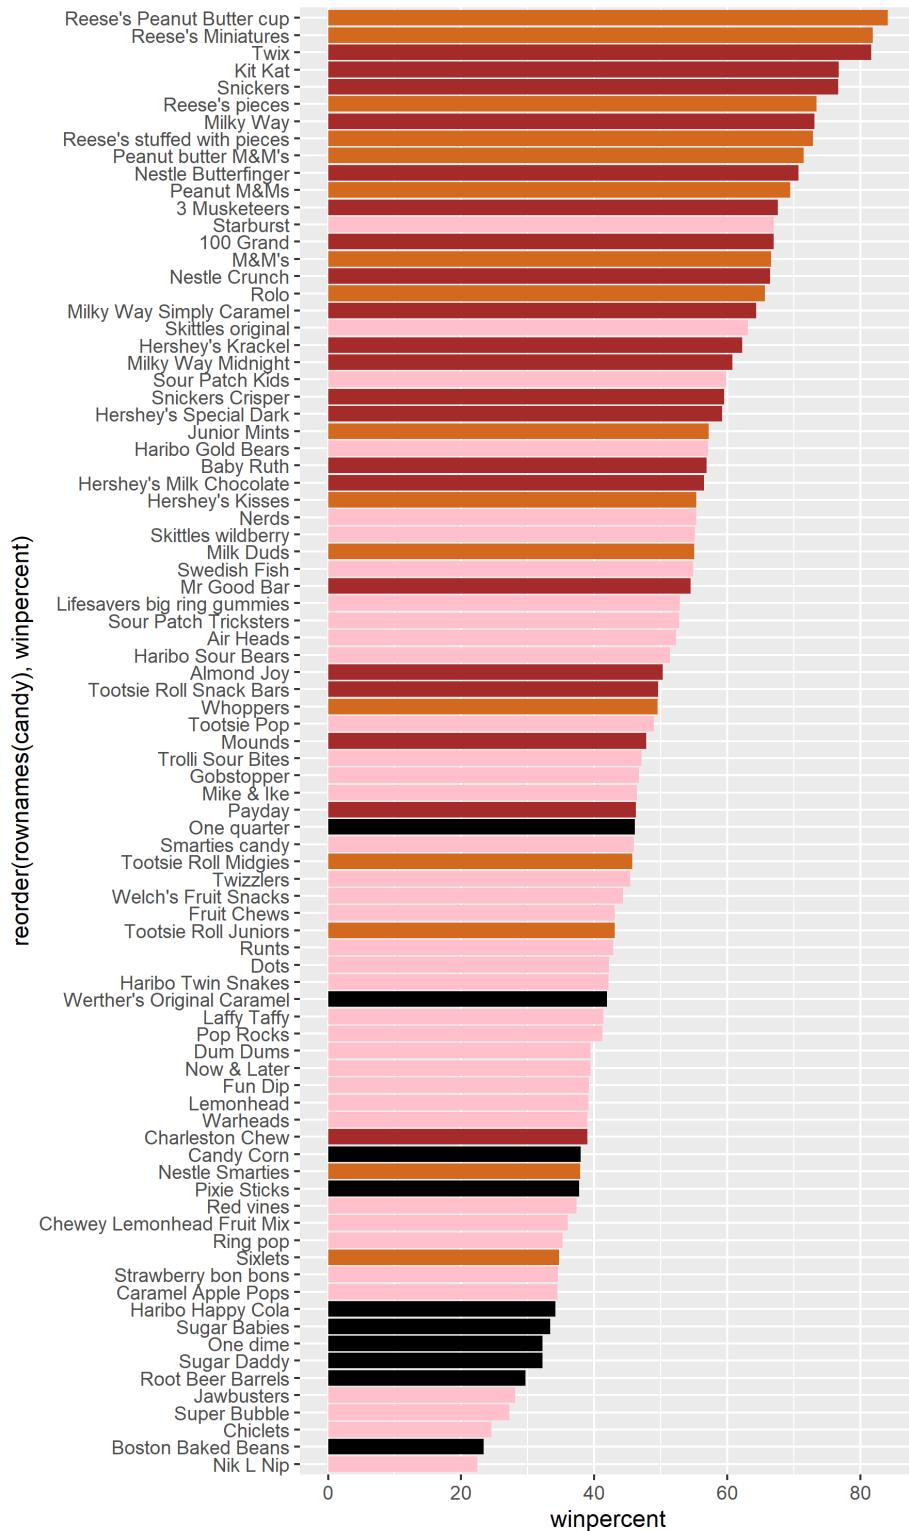
Q17. What is the worst ranked chocolate candy?

```
my_cols=rep("black", nrow(candy))
my_cols[as.logical(candy$chocolate)] = "chocolate"
my_cols[as.logical(candy$bar)] = "brown"
my_cols[as.logical(candy$fruity)] = "pink"
```

```
ggplot(candy) +
  aes(winpercent, reorder(rownames(candy), winpercent)) +
  geom_col(fill=my_cols)
```



```
ggsave("barplot3.png", height=10, width=6)
```



The worst ranked chocolate is sixlets

Q18. What is the best ranked fruity candy?

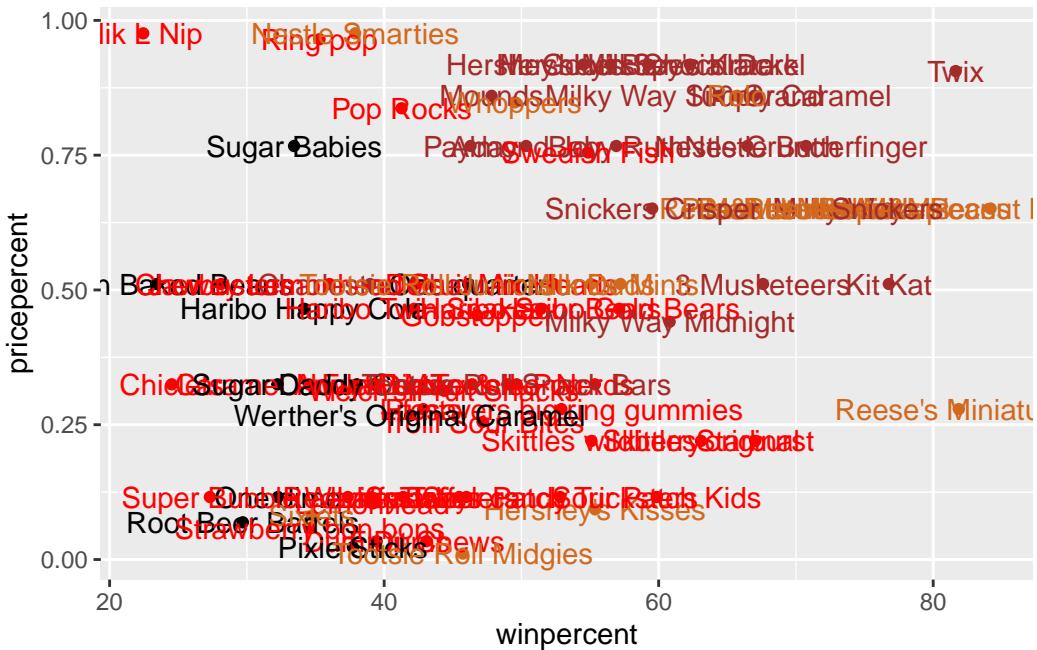
The best ranked fruity candy is starburst

Taking a look at pricepercent

Make a plot of winpercent vs the pricepercent

```
my_cols[as.logical(candy$fruity)] = "red"

ggplot(candy) +
  aes(x=winpercent, y=pricepercent, label=rownames(candy)) +
  geom_point(col=my_cols) +
  geom_text(col=my_cols)
```



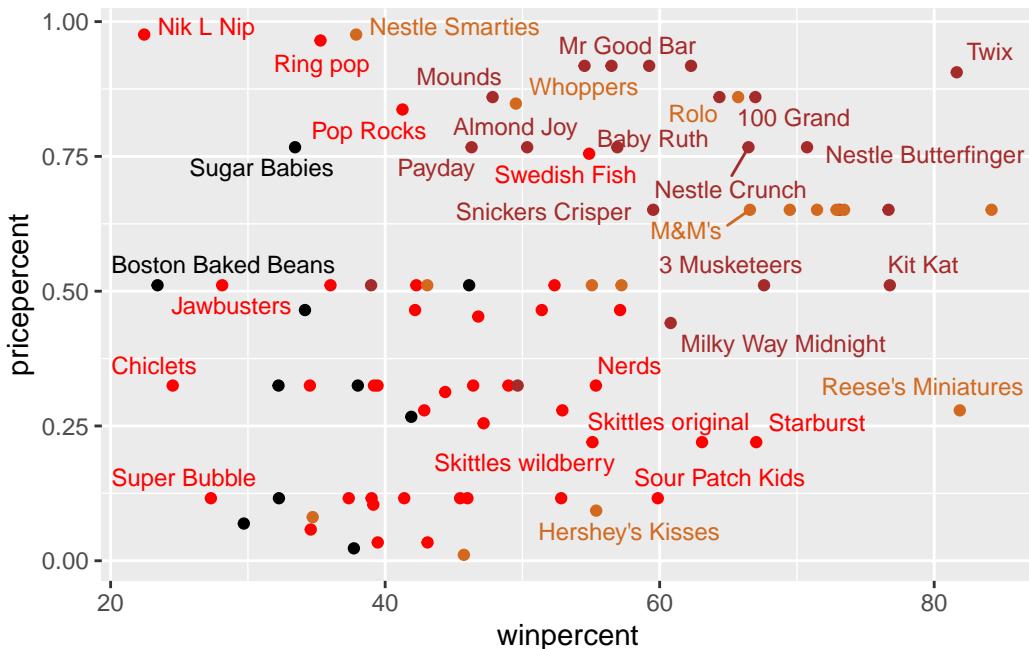
We can use **ggrepel** package for better label placement:

```
library(ggrepel)

my_cols[as.logical(candy$fruity)] = "red"
```

```
ggplot(candy) +
  aes(x=winpercent, y=pricepercent, label=rownames(candy)) +
  geom_point(col=my_cols) +
  geom_text_repel(col=my_cols, max.overlaps = 8, size = 3.3)
```

Warning: ggrepel: 52 unlabeled data points (too many overlaps). Consider increasing max.overlaps



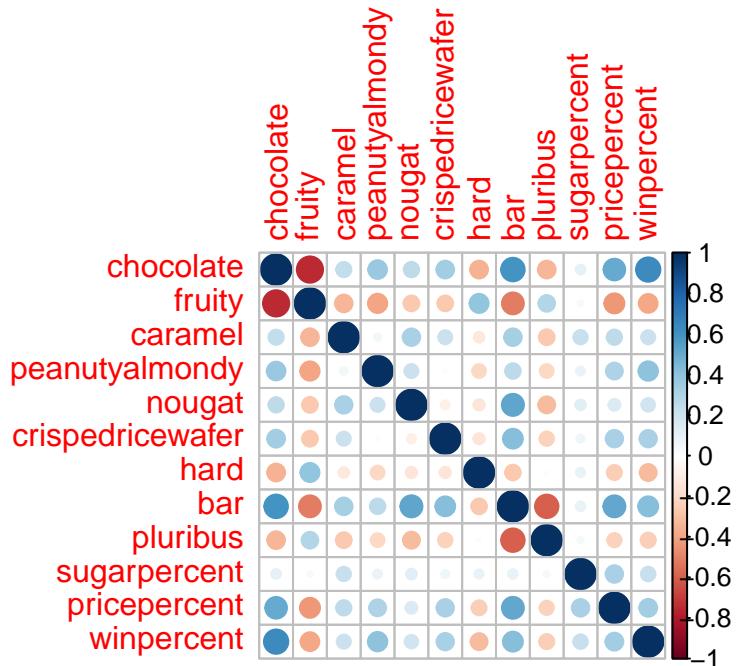
Exploring the correlation structure

Pearson correlation values range from -1 to +1

```
library(corrplot)
```

corrplot 0.95 loaded

```
cij <- cor(candy)
corrplot(cij)
```



Principal Component Analysis

```
pca <- prcomp(candy, scale=T)
summary(pca)
```

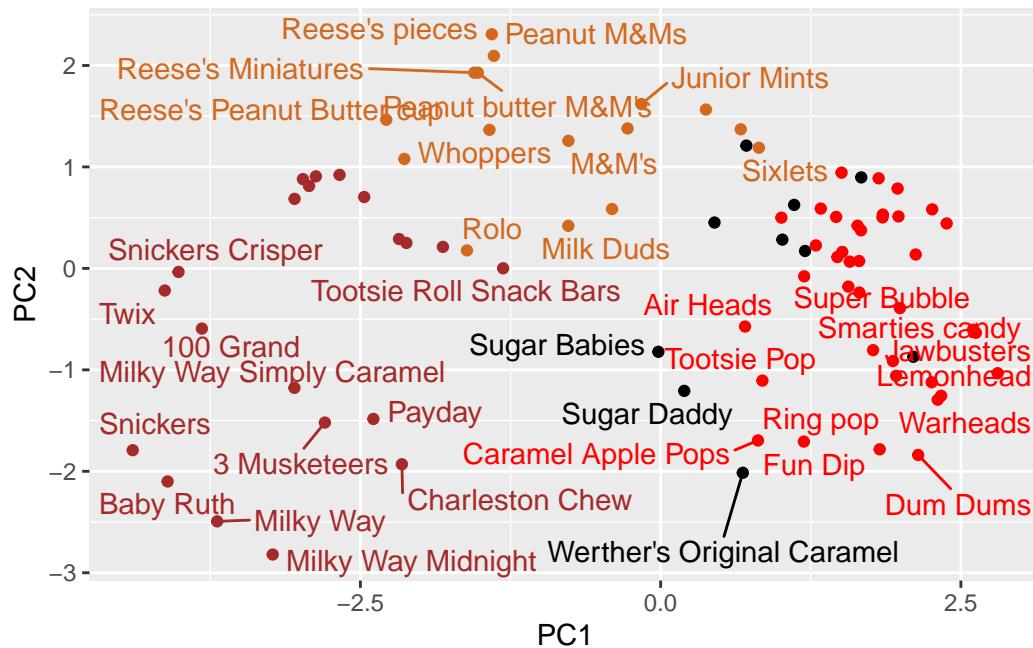
Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	2.0788	1.1378	1.1092	1.07533	0.9518	0.81923	0.81530
Proportion of Variance	0.3601	0.1079	0.1025	0.09636	0.0755	0.05593	0.05539
Cumulative Proportion	0.3601	0.4680	0.5705	0.66688	0.7424	0.79830	0.85369
	PC8	PC9	PC10	PC11	PC12		
Standard deviation	0.74530	0.67824	0.62349	0.43974	0.39760		
Proportion of Variance	0.04629	0.03833	0.03239	0.01611	0.01317		
Cumulative Proportion	0.89998	0.93832	0.97071	0.98683	1.00000		

The main results figure: the PCA score plot:

```
ggplot(pca$x) +
  aes(PC1, PC2, label=rownames(pca$x)) +
  geom_point(col=my_cols) +
  geom_text_repel(col=my_cols)
```

Warning: ggrepel: 48 unlabeled data points (too many overlaps). Consider increasing max.overlaps

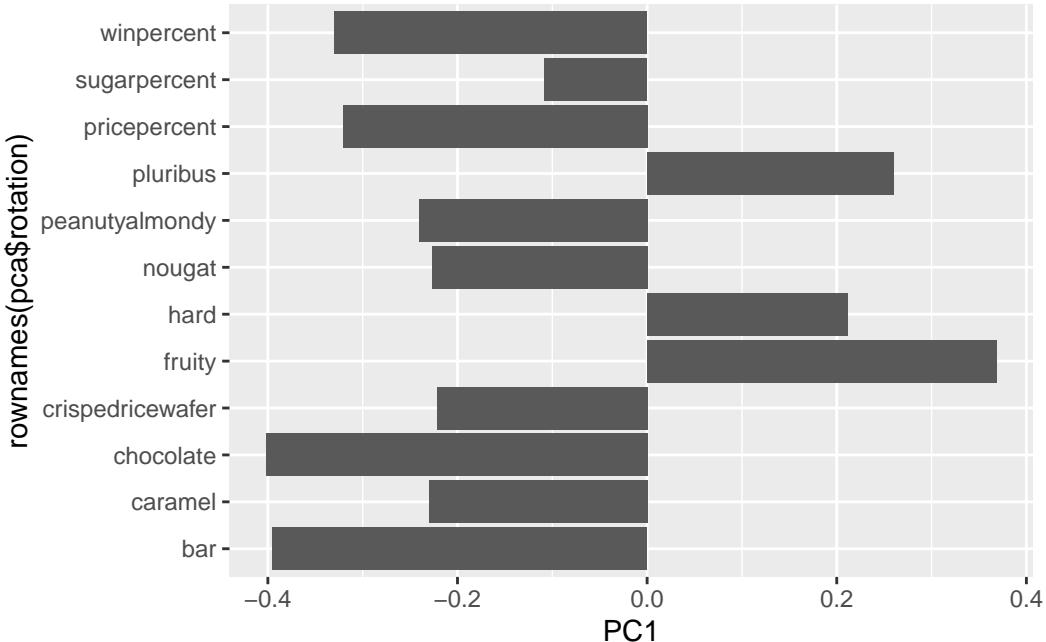


```
  labs(title="PCA Candy Space Map")
```

```
<ggplot2::labels> List of 1  
  $ title: chr "PCA Candy Space Map"
```

The “loadings” plot for PC1

```
ggplot(pca$rotation) +  
  aes(PC1, rownames(pca$rotation)) +  
  geom_col()
```



Q24. Complete the code to generate the loadings plot above. What original variables are picked up strongly by PC1 in the positive direction? Do these make sense to you? Where did you see this relationship highlighted previously?

Q25. Based on your exploratory analysis, correlation findings, and PCA results, what combination of characteristics appears to make a “winning” candy? How do these different analyses (visualization, correlation, PCA) support or complement each other in reaching this conclusion?